Lab 3
FPGAs as Accelerators

Introduction

The goal of this Lab is to implement the **LSAL algorithm** on the *x86, Arm* and *FPGA* of the *zedboard* and to *optimize* it in every case using **Valgrind**, **Intel advisor** and **Vitis HLS**.

LSAL stands for *local sequence alignment algorithm* and is used in *bioinformatics* and works by giving *an input of two strings, a database D of length M, and a query Q of length N.*
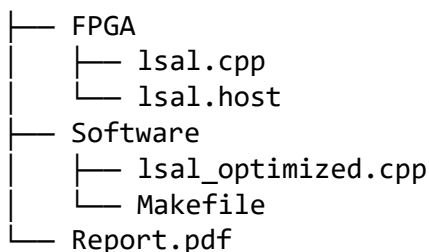
The *objective* is to find regions of similarity between subsequences of all possible lengths. An *M × N similarity matrix* S is constructed to hold similarity scores for each subsequence comparison.

A *second matrix* (called direction matrix) is formed to hold the direction path through matrix S to produce an *alignment*. The cell value of the direction matrix indicates the direction followed to reach that cell. The optimal alignment is produced by starting from the position of the maximum score in matrix S and following the directions found at the same index in the direction matrix until a zero in S is found.

In typical cases, the *size M* of the *database string* can be in the order of millions of *symbols*.

File Structure

The submitted files follow the following format.

```
├── FPGA
│   ├── lsal.cpp
│   └── lsal.host
├── Software
│   ├── lsal_optimized.cpp
│   └── Makefile
└── Report.pdf
```

The *FPGA folder* contains the *lsal_host.cpp* code that runs on the *ARM* of the *zedboard* and the *lsal.cpp* code that runs on the *FPGA*.

The *Software* folder contains the *lsal_optimized.cpp* code for the *software* and the *MAKEFILE* that allow it to compile *for the desired purpose*.

Software (x86 & Arm)

Workflow

The general workflow for the development of the code which is meant to run on x86 and ARM environments is as follows:

1. *First,* we **implement** the *LSAL algorithm without* considering the *CPU running time* and only taking into account the *functionality* of the code.
2. *Then* we **profile** the code using **Valgrind** and **Intel Advisor** while identifying the parts of the code that take the *most amount of time to run and which of them can change*.
3. *Depending* on the results of the previous step we change the necessary parts of the code but now *aiming for* **performance**.
4. *Lastly,* we **compare** the running times of the *optimized* and *unoptimized version* both by *CPU running times* and **Roofline analysis** *results*.

In both the unoptimized and optimized code we have added a **debug** *DEFINE* that allows us to debug any possible problem by just using the command *make debug* of our *MAKEFILE*.

Unoptimized

The **unoptimized code** represents an *implementation* of the *LSAL algorithm* where the output matrix is being calculated row-wise.

In the code we have added a **second for loop** *(compared to the lab input)* in order for the algorithm to *calculate* the values of the **first row** because it meets a lot of *edge cases* which would be *not efficient* to place it in the *main loop* as all those edge cases would need

to be checked for every value of the output arrays while this only needs to happen while the first row is being filled out.

In the main loop (that was already there from the lab input) we have added all the cases and the necessary calculations. *At the end of the for loop* we save the results both in the *similarity matrix* and the *direction matrix* while checking If the max index needs to be *"dethroned"*.

The results for a varying number of array sizes (Query and Database) are shown in the table below.

x86 vs ARM Unoptimized (sec)				
N	M	ARM	x86	Difference
32	32	0.000098	0.00002	79.5918%
32	65536	0.175177	0.037108	78.8169%
256	65536	1.402216	0.303335	78.3675%
256	300000	6.417753	1.388256	78.3685%

As we can see the *x86* is **78.82% faster** on *average* than the *ARM processor* of the *zedboard*.

Optimized

In order to profile the unoptimized code, we used **Valgrind** a *tool* that helps to identify parts of the code that take up most of the *execution time*.

Branching

The **branches** are some of the *instances* where the code was **slow** and *our solution* to the problem was to *reduce the if commands by* **unrolling** *the loops*.

For example, the **northwest** and the **west** during the calculation of cells of the *first column* **cannot be read** *because they are out of the bounds of the similarity array*. *Thus,* for the *first cell* of each *row* the calculation of *northwest and west* **is skipped** by unrolling the inner loop.

Division & Modulo

Also, calculations needed to find the *index* of the *similarity* and *direction array* are hard to calculate thus take **more time** to be executed (Mostly **modulo %** and **division /**). *So,* to *eliminate* the need to calculate the *x*

and y indexes of the matrices we **divide** the *main loop* to **two smaller** ones with one of them *representing the rows and the other the columns*.

Defines

It was noticed that **constant variables** were used to *define penalty, match, mismatch scores and enumerate the directions of the direction matrix (N, W, NW, C). That* may have a **bad effect** on performance because to access those values it is needed to *access memory*. So *even though the boost performance was minor* we transformed every constant variable to *defined values that became static during compilation (#define)*.

x86 vs ARM Optimized (sec)				
N	M	ARM	X86	Optimization
32	32	0.000045	0.000015	66.6667%
32	65536	0.076079	0.024057	68.3789%
256	65536	0.586488	0.196926	66.4228%
256	300000	2.685104	0.897239	66.5846%

Comparison

Lastly, we have run the **Roofline Analysis** that compares the *optimized* and *unoptimized code (when run on x86)* using the **Intel Advisor** as shown below:
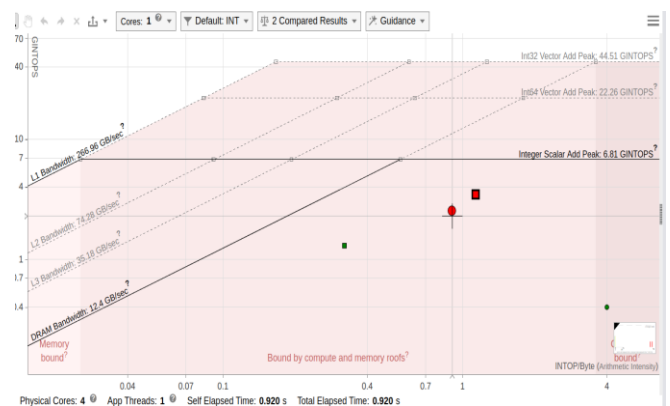

Image: Diagram of Roofline Model Analysis using Intel Advisor

The **red circle** represents the *unoptimized version* whereas the **red square** represents the *optimized version* of the code.

As we can see the *optimized version* has achieved *better parallelism* (as it is placed higher) and is even less dependent on the memory bandwidth (as it is placed further to the right).

With these changes the running times have dropped significantly with an average decrease of **32.65%** on the x86 and **56%** in the ARM processor of the zedboard.

As a conclusion, the optimized performs better on all aspects and has reduced running time both on x86 and on Arm by **44.32%** on *average*.

x86 (sec)				
N	M	Unoptimized	Optimised	Optimization
32	32	0.00002	0.000015	25.0000%
32	65536	0.037108	0.024057	35.1703%
256	65536	0.303335	0.196926	35.0797%
256	300000	1.388256	0.897239	35.3693%

ARM (sec)				
N	M	Unoptimized	Optimised	Optimization
32	32	0.000098	0.000045	54.0816%
32	65536	0.175177	0.076079	56.5702%
256	65536	1.402216	0.586488	58.1742%
256	300000	6.417753	2.685104	58.1613%

Hardware (FPGA)

During the last part of the lab, we use **Vitis HLS** to implement the *LSAL algorithm* on *FPGA* hardware.

In order to have the optimum outcome we defined a *streamlined* **workflow**:

1. *First* our team tries different **optimization methods** on *Visual Studio Code IDE*.
2. *Then functionality* of the *C code* is tested using *Vitis HLS's* **C Simulation**.
3. *At this point synthesis* and **Synthesis Design Report** of *Vitis HLS* is used to analyze the *resources in-use and the time estimations.*
4. *Next,* we use the provided *makefile* to **emulate** the *kernel design* using **QEMU**
5. *Last*ly, we test the design **natively** on the FPGA and *extract* the **timing** of *our implementation*

It is important to note that both *C Simulation, which tests functionality of the code before it gets synthesised,* and *software emulation using QEMU* are really important as a simple mistake can be detected and fixed long before the time it takes to have the kernel implemented into a bitstream.

Also, Synthesis Design Report is a great tool to have in order to make sure the HLS optimizations you use have been implemented in the correct way. *For example,* an easy mistake to make is placing a *#pragma* above a loop instead of inside it, which could result in vastly different implementations.

Algorithmic Optimizations

The main optimization introduced in the category of algorithmic optimizations is using **diagonals** to calculate the *similarity matrix*. As mentioned during the lecture, in this way we *limit data dependencies* and *"unlock"* **parallelism**.
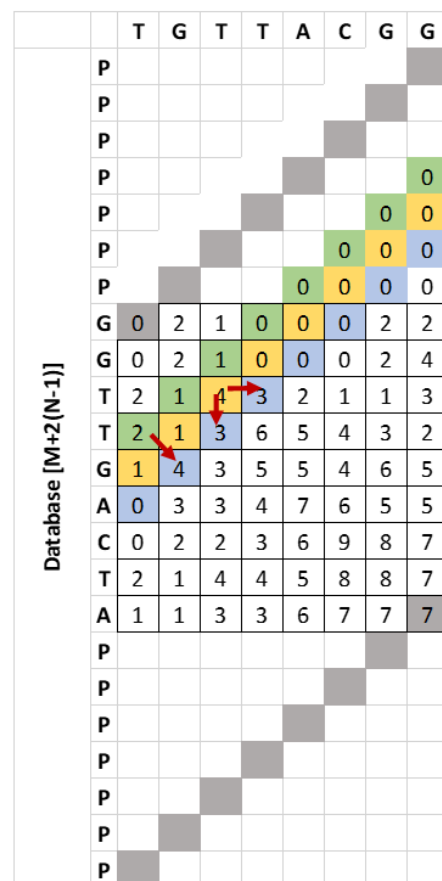

Figure 1: Showing diagonal parallelism, image provided during lecture.

This method requires that every diagonal is of the same size. Thus, we need to enlarge the database by $2(N-1)$, *where N is the size of query and M the size of database.*

To traverse the direction matrix **two nested loops**. *The outer one* counts the diagonals, *which are $M + (N -$*

1) *in crowd. The first diagonal* to be calculated is the *one that starts* at the *top right element* of the *extended direction matrix.*

The other one indicates which element of the diagonal is calculated, *from top right to bottom left. We decided to traverse* the diagonal in *such a way* because it *much easier* to calculate the *index variable* that is needed to store the *location* of the *maximum value.*

We realized that there are **data dependencies** on the *2 diagonals* above the one in calculation *during the calculation of a diagonal. That means* once the diagonal is calculated it can be moved to the buffer of the above diagonal.

Also, one thing noted to us is that at the beginning of each iteration of the loop we have a **write-after-read** on the *current diagonal, because* at the *end* of a *iteration* we write the diagonal we just calculated to main memory while at the *start* of the next we try to **write** to it. *That may be an* **obstacle** to having a **pipelined** *design* with **low II**.

To limit the **data dependencies,** we **double** the size of the *diagonal direction* and the *current diagonal (stores temporary similarity matrix values of the current diagonal)* matrices.

HLS Optimizations

This section includes every *optimization* that is **RTL specific** using *HLS.*

Array buffers

A method of making memory access faster is by using memory blocks that are near the processing unit of the acceleration. The designer in this case has 2 main options:

- by using **Block RAM** *(BRAM), that is a distinct memory block surrounded by Look-up Tables,* and
- by using the **Flip-Flops** inside the *Look-up Tables* of the *FGPA, which is faster but at the expense of making the rest of the LuT hard to utilize.*

Thus, for the best results a combination of the two is needed.

For our implementation the arrays have the following characteristics:

Array	Description	#pragma	Definition
string1	Buffer of query	PARTITION	complete
string2	Buffer of database	PARTITION	factor=2 cyclic
current_diag	Buffer of current diagonal	PARTITION	complete
up_diag	Buffer of the above diagonal	PARTITION	complete
upper_diag	Buffer of 2 diagonals above the current one	PARTITION	complete
direction_buf	Buffer of current diagonal	PARTITION	factor=32 cyclic

The preferred definition is complete because It transforms all the values of the array into separate registers which makes data transfer as fast as possible.

Now for the big arrays of string2 and direction_buf we did not use the complete definition because there are not enough registers and we used cyclic BRAM because in instances where consecutive values of an array are read, we don't want to be limited by the read ports of the Bram. So, by using cyclic BRAM'S with a factor of 32 we essentially have 32 Brams (where each one of them has a read port) with all the different values that we need which unlocks further parallelism.

Arbitrary Elements

Another easy **optimization** to make is to *reduce the* **size of the bus** used on the *design.*

For example, in the case of the **query** and **database** we know *as a fact* that they need to *enumerate* just the *4 nucleotides* in existence plus *one* more *undefined state we use in combination with the diagonals optimization.* Thus, we only need to use **3-bits** to those **enumerate 5 states**,
$(2^3 = 8 > 5\ states)$.

In the exact same fashion, the bits used to represent the **direction matrix** are reduced to **3-bits**.

To be noted other data structures such as the *similarity matrix*, it's *diagonal buffers* and *max_index* **must not** be represented by a ***smaller number of bits*** because the *RTL designers* cannot *easily* calculate *the maximum value* each structure is going to store.

Comparison

		FPGA			
N	M	Unoptimized	Sol1	Sol2	Difference
32	32	0.002565564	0.000817	-	68.1421%
32	65536	3.268688927	0.380162	0.17579	94.6219%
256	65536	5.5359965	2.176	-	60.6936%

In the first solution of the FPGA code the only "optimization" are the buffers and we have not yet achieved to pipeline the compute matrice's main for loop and thus the time is still very high and away from the target of 10 ms.

In the second solution we have focused on the inputs of 32 65536 and we have achieved almost half the time because now we have a pipeline with II=32 which is not ideal, but it is better than anything. Also, in this iteration we have corrected some write after read dependencies that we mentioned above and made parallelism a particularly difficult task.

Results

FPGA vs CPU

Finally as a result we have the following:

		FPGA vs x86	
N	M	x86	FPGA
32	32	0.00002	0.0008173
32	65536	0.037108	0.1757924
256	65536	0.303335	2.176

As you can see the results are not the expected and we have not managed to reach the goal of 10 ms for the FPGA for the inputs of 32 65536. We have though achieved the goal of <30 ms for the x86 which is still a success.